Lecture 10 Big Data Storage Technologies

Dr. Lydia Wahid

Agenda

- **Introduction**
- On-Disk Storage:
- Distributed File System
- RDBMS Databases
- NoSQL Databases
- NewSQL Databases
- In-Memory Storage:
- In-Memory Data Grids
- In-Memory Database (Relational, NoSQL, NewSQL)



Introduction

Introduction

- >Storage technology has continued to evolve over time.
- The need to **store Big Data** has radically altered the relational, database-centric view that has been embraced by Enterprise ICT.
- The bottom line is that relational technology is simply **not scalable** in a manner to support **Big Data volumes**.
- Also, businesses can find genuine value in processing semi-structured and unstructured data, which are generally incompatible with relational approaches.

Introduction

- ➤ Big Data has pushed the storage boundary to unified views of the available memory and **disk storage** of a cluster.
- If more storage is needed, horizontal scalability allows the expansion of the cluster through the addition of **more nodes**.
- Innovative approaches deliver **realtime** analytics via **in-memory storage**.
- Even batch-based processing is accelerated by the performance of Solid State Drives (SSDs), which have become less expensive.



On-Disk Storage

On-Disk Storage

On-disk storage generally utilizes **low cost** hard-disk drives for long-term storage.

➤On-disk storage can be implemented via a **distributed file system** or a **database**.



On-Disk Storage: Distributed File Systems

On-Disk Storage: Distributed File Systems

- Distributed file systems, like any file system, support schema-less data storage.
- In general, a distributed file system storage device provides **redundancy** and high **availability** by copying data to multiple locations via **replication**.
- A storage device that is implemented with a distributed file system provides *simple*, *fast access* data storage that is capable of storing large datasets that are non-relational in nature, such as **semi-structured** and **unstructured** data.

On-Disk Storage: Distributed File Systems

- Although based on straightforward file **locking** mechanisms for concurrency control, it provides **fast read/write** capability, which addresses the **velocity** characteristic of Big Data.
- Distributed file system is **not ideal** for datasets comprising a **large number of small files** as this creates excessive **disk-seek activity**, **slowing down** the overall data access.
- Due to these limitations, distributed file systems work best with **fewer** but larger files accessed in a sequential manner.

On-Disk Storage: Distributed File Systems

- Multiple smaller files are generally **combined** into a single file to enable optimum *storage* and *processing*.
- This allows the distributed file systems to have increased performance when data must be accessed in **streaming mode** with **no random** reads and writes.
- A distributed file system storage device is suitable when large datasets of raw data are to be stored or when archiving of datasets is required.
- It should be noted that distributed file systems do not provide the ability to search the contents of files



- Relational database management systems (RDBMSs) are good for handling transactional workloads involving small amounts of data with random read/write properties.
- ➤ RDBMSs are **ACID-compliant**, and they are generally restricted to a **single node**. For this reason, RDBMSs **do not provide** *redundancy* and *fault tolerance*.
- To handle large volumes of data arriving at a fast pace, relational databases generally **need to scale**. RDBMSs employ **vertical scaling**, not horizontal scaling, which is a *more costly* and *disruptive* scaling strategy.

- Note that some relational databases, for example IBM DB2 pureScale, Sybase ASE Cluster Edition, Oracle Real Application Clusters (RAC) and Microsoft Parallel Data Warehouse (PDW), are capable of being run on clusters.
- However, these database clusters still **use shared storage** that can act as a single point of failure.
- Relational databases need to be manually sharded, mostly using application logic.
- This means that the application logic needs to know which shard to query in order to get the required data. This further **complicates** data processing when data from **multiple shards** is required.

- ➤ Consider the following scenario:
 - 1. A user writes a record (id = 2).
 - 2. The **application logic** determines **which shard** it should be written to.
 - 3. It is **sent to the shard** determined by the application logic.
 - 4. The user reads a record (id = 4), and the application logic determines which shard contains the data.
 - 5. The data is read and returned to the application.
 - 6. The application then returns the record to the user.

- Relational databases generally require data to **adhere to a schema**. As a result, storage of **semi-structured** and **unstructured** data whose schemas are non-relational is **not directly supported**.
- Furthermore, with a relational database schema conformance is validated at the time of data insert or update by checking the data against the constraints of the schema. This introduces overhead that creates latency.
- This latency makes relational databases a less than ideal choice for storing **high velocity** data that needs a highly available database storage device with *fast* data *write* capability.



On-Disk Storage: NoSQL Databases

- Not-only SQL (NoSQL) refers to technologies used to develop next generation non-relational databases that are highly scalable and fault-tolerant.
- ➤ Below is a list of the principal characteristics of NoSQL storage devices that differentiate them from traditional RDBMSs:
 - Schema-less data model Data can exist in its raw form.
 - **Highly available** This is built on cluster-based technologies that provide fault.
 - Eventual consistency Data reads across multiple nodes but may not be consistent immediately after a write. However, all nodes will eventually be in a consistent state.

- ➤ Below is a list of the principal characteristics of NoSQL storage devices that differentiate them from traditional RDBMSs (cont.):
 - Scale out rather than scale up More nodes can be added to obtain additional storage with a NoSQL database.
 - **BASE**, **not ACID** BASE compliance requires a database to maintain high *availability* in the event of network/node failure, while not requiring the database to be in a *consistent* state whenever an update occurs. NoSQL storage devices are generally *AP* or *CP*.
 - **API driven data access** Data access is generally supported via API based queries, including *RESTful APIs*, whereas some implementations may also provide *SQL-like query* capability.

- ➤ Below is a list of the principal characteristics of NoSQL storage devices that differentiate them from traditional RDBMSs (cont.):
 - **Auto sharding and replication** To support horizontal scaling and provide high availability, a NoSQL storage device automatically employs sharding and replication techniques.
 - **Distributed query support** NoSQL storage devices maintain consistent query behavior across multiple shards.
 - **Polyglot persistence** The use of NoSQL storage does not mandate retiring traditional RDBMSs. In fact, both can be used at the same time, thereby supporting polyglot persistence, which is an approach of *persisting data using different types of storage technologies within the same solution architecture*. This is good for developing systems requiring structured as well as semi/unstructured data.

- ➤ Below is a list of the principal characteristics of NoSQL storage devices that differentiate them from traditional RDBMSs (cont.):
 - **Aggregate-focused** Unlike relational databases that are most effective with fully normalized data, NoSQL storage devices store denormalized aggregated data (an entity containing merged, often nested, data for an object) thereby eliminating the need for joins and extensive mapping between application objects and the data stored in the database. One exception, however, is that graph database storage devices (introduced shortly) are not aggregate-focused.

NoSQL Databases: Rational

The emergence of NoSQL storage devices can primarily be attributed to the *volume*, *velocity* and *variety* characteristics of Big Data datasets:

• Volume:

- The storage requirement of ever increasing data volumes commands the use of databases that are **highly scalable** while keeping costs down for the business to remain competitive.
- NoSQL storage devices fulfill this requirement by providing **scale out** capability while using **inexpensive** commodity servers.

NoSQL Databases: Rational

The emergence of NoSQL storage devices can primarily be attributed to the *volume*, *velocity* and *variety* characteristics of Big Data datasets:

• Velocity:

- The fast influx of data requires databases with fast access data write capability.
- NoSQL storage devices enable fast writes by using **schema-on-read** rather than **schema-on-write** principle.
- Being highly available, NoSQL storage devices can ensure that write latency does not occur because of node or network failure.

NoSQL Databases: Rational

The emergence of NoSQL storage devices can primarily be attributed to the *volume*, *velocity* and *variety* characteristics of Big Data datasets (cont.):

• Variety:

- A storage device needs to handle **different data formats** including *documents*, *emails*, *images* and *videos* and *incomplete* data.
- NoSQL storage devices can store these different forms of **semi-structured** and **unstructured** data formats.
- At the same time, NoSQL storage devices are able to store schema-less data and incomplete data with the added ability of making schema changes as the data model of the datasets evolve.
- In other words, NoSQL databases support schema evolution.

NoSQL Databases: Types

- NoSQL storage devices can mainly be divided into four types based on the way they store data:
 - Key-value
 - Document
 - Column-family
 - Graph

NoSQL Databases: Types



key	value
631	John Smith, 10.0.30.25, Good customer service
365	1001010111011011110111010101010101010110011010
198	<customerid>32195</customerid> <total>43.25</total>

An example of key-value NoSQL storage.

An example of document NoSQL storage.

studentld	personal details	address	modules history
821	FirstName: Cristie LastName: Augustin DoB: 03-15-1992 Gender: Female Ethnicity: French	Street: 123 New Ave City: Portland State: Oregon ZipCode: 12345 Country: USA	Taken: 5 Passed: 4 Failed: 1
742	FirstName: Carlos LastName: Rodriguez MiddleName: Jose Gender: Male	Street: 456 Old Ave City: Los Angeles Country: USA	Taken: 7 Passed: 5 Failed: 2

An example of column-family NoSQL storage.

TV

Brand: Sony

Price: \$1,200 Size: 55"

7

bought

Customer

Name: James

Age: 56

Gender: Male

Customer

bought

Name: Claire

Age: 23

Gender: Female

Customer

Name: Josh

Age: 35

viewed

Gender: Male

An example of graph NoSQL storage.

- > Key-value storage devices store data as key-value pairs and act like hash tables.
- The table is a list of values where each value is identified by a key. The value is **opaque** to the database and is typically stored as a **BLOB** (Binary large Objects).
- The value stored can be any aggregate, ranging from sensor data to videos.
- ➤ Value look-up can only be performed via the keys as the database is oblivious to the details of the stored aggregate. Partial updates are not possible. An update is either a **delete** or an **insert** operation.
- Examples of key-value storage devices include Riak, Redis, and Amazon Dynamo DB.

- ➤ Key-value writes are quite **fast**. Based on a simple storage model, key-value storage devices are **highly scalable**.
- keys are the only means of retrieving the data.
- A single collection can hold multiple data formats.

key	value	
631	John Smith, 10.0.30.25, Good customer service	<text< td=""></text<>
365	10101101010110101011010101010101010111010	✓ image
198	<customerid>32195</customerid> <total>43.25</total>	← XML

- A key-value storage device is **appropriate** when:
 - unstructured data storage is required
 - high performance read/writes are required
 - the value is fully identifiable via the key alone
 - value is a standalone entity that is not dependent on other values
 - values have a comparatively simple structure or are binary
 - query patterns are simple, involving insert, select and delete operations only
 - stored values are manipulated at the application layer

- A key-value storage device is **inappropriate** when:
 - applications require searching or filtering data using attributes of the stored value
 - relationships exist between different key-value entries
 - a group of keys' values need to be updated in a single transaction
 - multiple keys require manipulation in a single operation
 - schema consistency across different values is required
 - update to individual attributes of the value is required

- Document storage devices also store data as key-value pairs.
- However, unlike key-value storage devices, the stored value is a document that can be **queried** by the database.
- These documents can have a complex nested structure.
- Like key-value storage devices, most document storage devices provide collections or buckets (like tables) into which key-value pairs can be **organized**.
- The documents can be encoded using either a text-based encoding scheme, such as **XML** or **JSON**, or using a **binary** encoding scheme.
- Examples of document storage devices include MongoDB, CouchDB, and Terrastore.

- The main differences between document storage devices and key-value storage devices are as follows:
 - document storage devices are value-aware
 - the stored value is self-describing; the schema can be inferred from the structure of the value or a reference to the schema for the document is included in the value
 - a select operation can reference a field inside the aggregate value
 - a select operation can retrieve a part of the aggregate value
 - partial updates are supported; therefore a subset of the aggregate can be updated

- Each document can have a **different schema**; therefore, it is possible to store different types of documents in the same collection or bucket.
- Additional fields can be added to a document after the initial insert, thereby providing flexible schema support.

```
{
    invoiceld:37235,
    date:19600801,
    custld:29317,
    items:[
        {itemId:473,quantity:2},
        {itemId:971,quantity:5}
    ]
}
```

An example of document NoSQL storage.

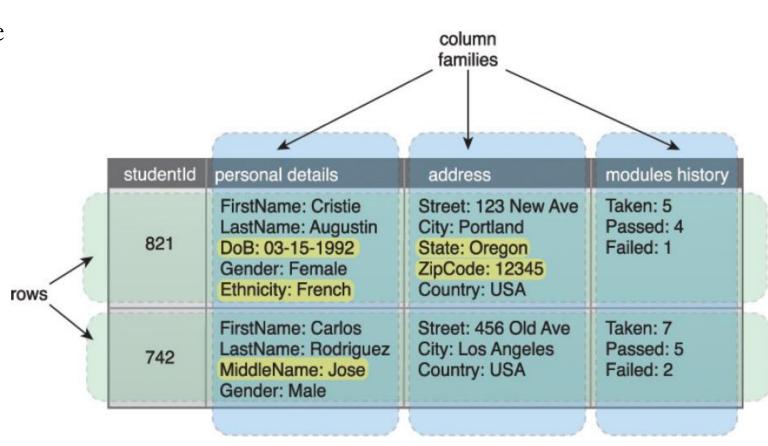
- A document storage device is **appropriate** when:
 - storing semi-structured document-oriented data comprising flat or nested schema
 - schema evolution is a requirement as the structure of the document is either unknown or is likely to change
 - applications require a partial update of the aggregate stored as a document
 - searches need to be performed on different fields of the documents
 - query patterns involve insert, select, update and delete operations

NoSQL Databases: Document

- A document storage device is **inappropriate** when:
 - multiple documents need to be updated as part of a single transaction
 - performing operations that need joins between multiple documents
 - schema enforcement for achieving consistent query design is required as the document structure may change between successive query runs, which will require restructuring the query
 - the stored value is not self-describing and does not have a reference to a schema

- Column-family storage devices store data much like a traditional RDBMS but group related columns together in a row, resulting in column-families.
- Each column can be a collection of related columns itself, referred to as a super-column.
- Each super-column can contain an arbitrary number of related columns that are generally retrieved or updated as a single unit.
- Each row consists of multiple column-families and can have a different set of columns, thereby manifesting flexible schema support.
- Each row is identified by a **row key**.
- Examples of column-family storage devices include Cassandra, HBase and Amazon SimpleDB.

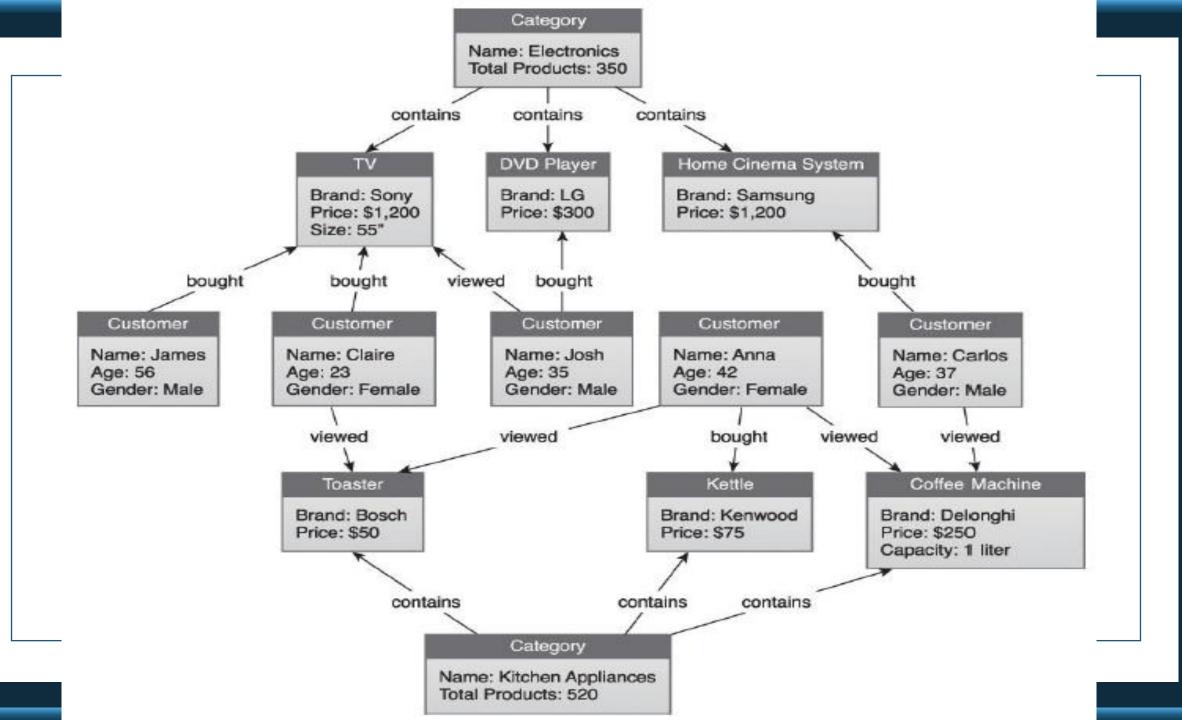
- Column-family storage devices provide fast data access with random read/write capability.
- They store different column-families in separate physical files, which improves query responsiveness as only the required column-families are searched.



- A column-family storage device is **appropriate** when:
 - data represents a tabular structure, each row consists of a large number of columns and nested groups of interrelated data exist
 - support for schema evolution is required as column families can be added or removed without any system downtime
 - certain fields are mostly accessed together, and searches need to be performed using field values
 - query patterns involve insert, select, update and delete operations

- A column-family storage device is **inappropriate** when:
 - relational data access is required; for example, joins
 - ACID transactional support is required
 - SQL-compliant queries need to be executed
 - query patterns are likely to change frequently because that could initiate a corresponding restructuring of how column-families are arranged

- >Graph storage devices are used to persist inter-connected entities.
- ➤ Unlike other NoSQL storage devices, where the emphasis is on the structure of the entities, graph storage devices place emphasis on storing the linkages between entities.
- Entities are stored as nodes and are also called vertices, while the linkages between entities are stored as edges.
- In RDBMS parlance, each node can be thought of a **single row** while the edge denotes a **join**.
- Examples include Neo4J, Infinite Graph and OrientDB.



- Nodes can have more than one type of link between them through multiple edges.
- Each node can have attribute data as key-value pairs, such as a customer node with ID, name and age attributes.
- Each edge can have its own attribute data as key-value pairs, which can be used to further filter query results.
- Having multiple edges are similar to defining multiple **foreign keys** in an RDBMS.

- Generally, graph storage devices provide consistency via **ACID** compliance.
- The degree of usefulness of a graph storage device depends on the **number** and **types** of edges defined between the nodes.
- The greater the number and more diverse the edges are, the more diverse the types of queries it can handle.
- Graph storage devices generally allow adding new types of nodes without making changes to the database.
- This also enables defining additional links between nodes as new types of relationships or nodes appear in the database.

- A graph storage device is **appropriate** when:
 - interconnected entities need to be stored
 - querying entities based on the type of relationship with each other rather than the attributes of the entities
 - finding groups of interconnected entities
 - finding distances between entities in terms of the node traversal distance
 - mining data with a view toward finding patterns

- A graph storage device is **inappropriate** when:
 - updates are required to a large number of node attributes or edge attributes, as this involves searching for nodes or edges, which is a costly operation
 - entities have a large number of attributes or nested data—it is best to store lightweight entities in a graph storage device while storing the rest of the attribute data in a separate non-graph NoSQL storage device
 - queries based on the selection of node/edge attributes dominate node traversal queries

- As mentioned **MongoDB** is a document-based NOSQL system.
- Individual documents are stored in a collection.
- The operation **createCollection** is used to create each collection.
- For example, the following command can be used to create a collection called **project** to hold PROJECT objects from the COMPANY database:
 - db.createCollection("project", { capped: true, size: 1310720, max: 500 })
 - The first parameter "project" is the **name** of the collection, which is followed by an optional document that specifies **collection options**. In our example, the collection is **capped**; this means it has upper limits on its storage space (**size**) and number of documents (**max**).

- For our example, we will create another document collection called **worker** to hold information about the EMPLOYEEs who work on each project; for example:
 - db.createCollection("worker", { capped: true, size: 5242880, max: 2000 }))
- Each document in a collection has a unique **ObjectId** field, called **_id**, which is automatically indexed in the collection unless the user explicitly requests no index for the **_id** field.

► MongoDB CRUD Operations:

• MongoDb has several **CRUD operations**, where CRUD stands for (create, read, update, delete). Documents can be *created* and inserted into their collections using the **insert** operation, whose format is:

db.<collection_name>.insert(<document(s)>)

• The parameters of the insert operation can include either a single document or an array of documents, as shown in the following figure:

```
inserting the documents into their collections "project" and "worker":

db.project.insert( { _id: "P1", Pname: "ProductX", Plocation: "Bellaire" } )

db.worker.insert( [ { _id: "W1", Ename: "John Smith", ProjectId: "P1", Hours: 32.5 }.

{ _id: "W2", Ename: "Joyce English", ProjectId: "P1",

Hours: 20.0 } ] )
```

► MongoDB CRUD Operations:

• The delete operation is called *remove*, and the format is:

db.<collection_name>.remove(<condition>)

The documents to be removed from the collection are specified by a Boolean condition on some of the fields in the collection documents.

- There is also an **update** operation, which has a condition to select certain documents, and a *\$set* clause to specify the update.
- For read queries, the main command is called *find*, and the format is:

db.<collection_name>.find(<condition>)



On-Disk Storage: NewSQL Databases

NewSQL Databases

- NoSQL storage devices are **highly scalable**, **available**, **fault-tolerant** and **fast** for read/write operations.
- However, they do not provide the same transaction and consistency support as exhibited by **ACID** compliant RDBMSs. Following the BASE model, NoSQL storage devices provide **eventual consistency** rather than **immediate consistency**.
- They therefore will be in a soft state while reaching the state of eventual consistency.
- As a result, they are **not appropriate** for use when **implementing large** scale transactional systems.

NewSQL Databases

NewSQL storage devices combine the **ACID** properties of **RDBMS** with the **scalability and fault tolerance** offered by **NoSQL** storage devices.

NewSQL databases generally **support SQL** compliant syntax for data definition and data manipulation operations, and they often use a logical **relational data model** for data storage.

NewSQL Databases

- NewSQL databases can be used for developing **OLTP** (online transaction processing) systems with very high volumes of transactions, for example a banking system.
- They can also be used for **realtime analytics** as some implementations leverage **in-memory storage**.

Examples of NewSQL databases include VoltDB, NuoDB and InnoDB.



- This section presents in-memory storage as a means of providing options for highly performant, and advanced data storage.
- An in-memory storage device generally utilizes **RAM**, the main memory of a computer, as its storage medium to provide **fast data access**.

The growing capacity and decreasing cost of RAM, coupled with the increasing read/write speed of solid state hard drives, has made it possible to develop in-memory data storage solutions.

- Storage of data in memory eliminates the latency of disk I/O and the data transfer time between the main memory and the hard drive.
- This overall reduction in data read/write latency makes data processing much faster.
- In-memory storage device capacity can be increased massively by horizontally scaling the cluster that is hosting the in-memory storage device.
- Cluster-based memory enables storage of large amounts of data, including Big Data datasets, which can be accessed **considerably faster** when compared with an on-disk storage device.

- In-memory storage significantly reduces the overall execution time of Big Data analytics, thus enabling real-time Big Data analytics.
- An example of a sequential read of **1 MB** of data from an in-memory storage device takes around **0.25 ms**. The same amount of data from an on-disk storage device takes around **20 ms**.

This demonstrates that reading data from in-memory storage is approximately **80 times** faster than on-disk storage.

- Primarily, in-memory storage enables making sense of the fast influx of data in a Big Data environment (velocity characteristic) by providing a storage medium that facilitates **real-time insight generation**.
- This supports making quick business decisions for mitigating a threat or taking advantage of an opportunity.
- A Big Data in-memory storage device is implemented over a cluster, providing **high availability** and **redundancy**. Therefore, horizontal scalability can be achieved by simply adding more nodes or memory.

- When compared with an on-disk storage device, an in-memory storage device is **expensive** because of the higher cost of memory as compared to a disk-based storage device.
- Apart from being expensive, in-memory storage devices do not provide the same level of support for **durable data storage**. The price factor further affects the achievable **capacity** of an in-memory device when compared with an on-disk storage device.

- An in-memory storage device is **appropriate** when:
 - data arrives at a **fast pace** and **requires real-time analytics** or **event stream processing**
 - continuous or always-on analytics is required, such as operational BI and operational analytics
 - interactive query processing and real-time data visualization needs to be performed
 - developing low latency Big Data solutions

- An in-memory storage device is **inappropriate** when::
 - data processing consists of batch processing
 - very large amounts of data need to be persisted in-memory for a long time in order to perform in-depth data analysis
 - performing strategic BI or strategic analytics that involves access to **very** large amounts of data and involves batch data processing
 - datasets are extremely large and do not fit into the available memory
 - an enterprise has a **limited budget**, as setting up an in-memory storage device may require upgrading nodes, which could either be done by node replacement or by adding more RAM

- In-memory storage devices can be implemented as:
 - In-Memory Data Grid (IMDG)
 - In-Memory Database (IMDB)

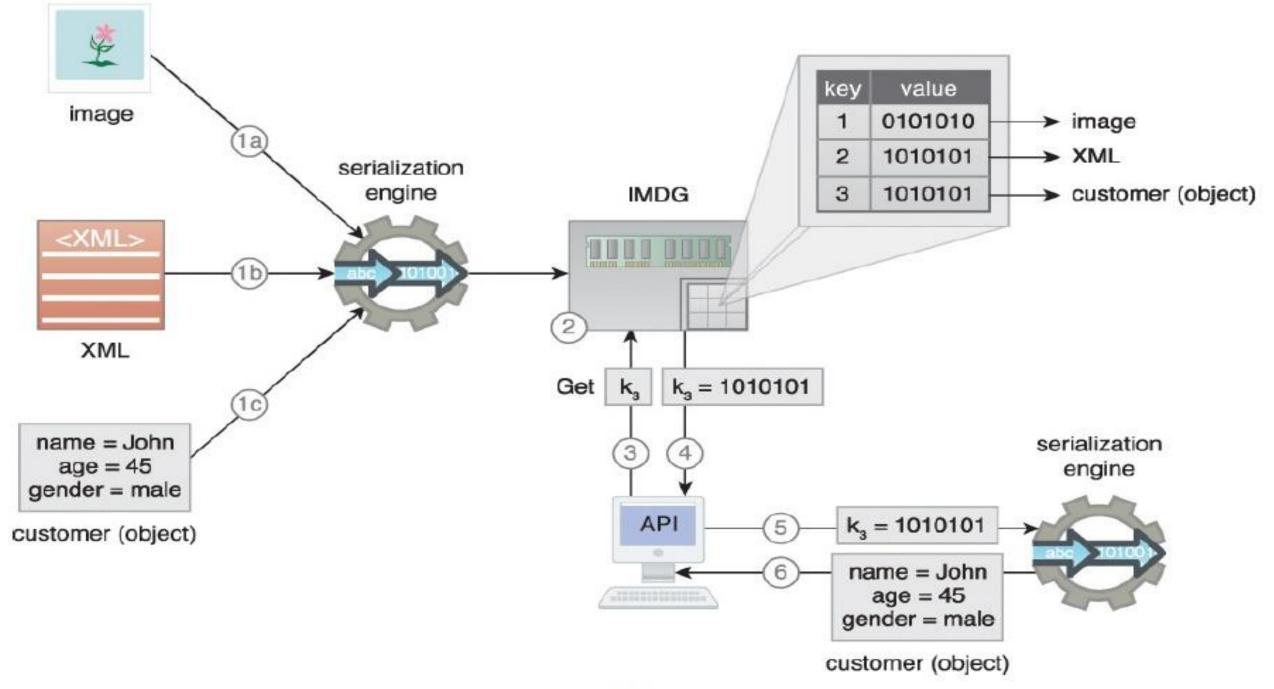
Although both of these technologies use memory as their underlying data storage medium, what makes them distinct is **the way data is stored in the memory**.



In-Memory Storage:
In-Memory Data Grids

IMDGs store data in memory as **key-value pairs** across multiple nodes where the keys and values can be any business object or application data in serialized form.

This supports schema-less data storage through storage of semi/unstructured data. Data access is typically provided via APIs.



An IMDG storage device.

- The previous figure describes the following scenario:
 - 1. An image (a), XML data (b) and a customer object (c) are first serialized using a serialization engine.
 - 2. They are then stored as key-value pairs in an IMDG.
 - 3. A client requests the customer object via its key.
 - 4. The value is then returned by the IMDG in serialized form.
 - 5. The client then utilizes a serialization engine to deserialize the value to obtain the customer object...
 - 6. ... in order to manipulate the customer object.

Nodes in IMDGs keep themselves synchronized and collectively provide high availability, fault tolerance and consistency.

In comparison to **NoSQL**'s eventual consistency approach, IMDGs support **immediate consistency**.

As compared to relational **IMDBs** (discussed later), IMDGs provide **faster** data access as IMDGs store non-relational data as objects.

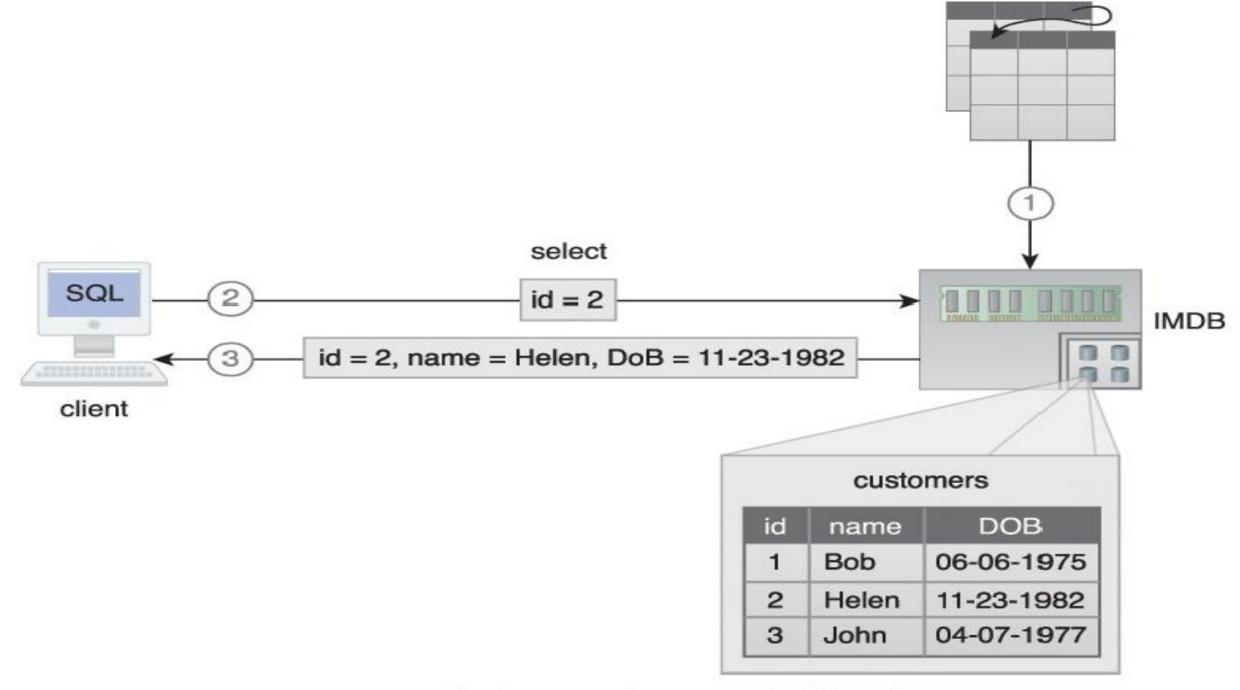
- ➤ Realtime processing engines can make use of IMDG where high velocity data is stored in the IMDG as it arrives and is processed there before being saved to an on-disk storage device, or data from the on-disk storage device is copied to the IMDG.
- This makes data processing orders of magnitude faster.
- ➤IMDGs may also support <u>in-memory MapReduce</u> that helps to reduce the latency of disk based MapReduce processing.

- IMDGs can be added to existing Big Data solutions by introducing them between the existing **on-disk storage device** and the **data processing** application.
- Note that some IMDG implementations may also provide limited or full **SQL support**.
- In a Big Data solution environment, IMDGs are often deployed together with **on-disk storage** devices that act as the **backend storage**.
- Examples include In-Memory Data Fabric, Hazelcast and Oracle Coherence.



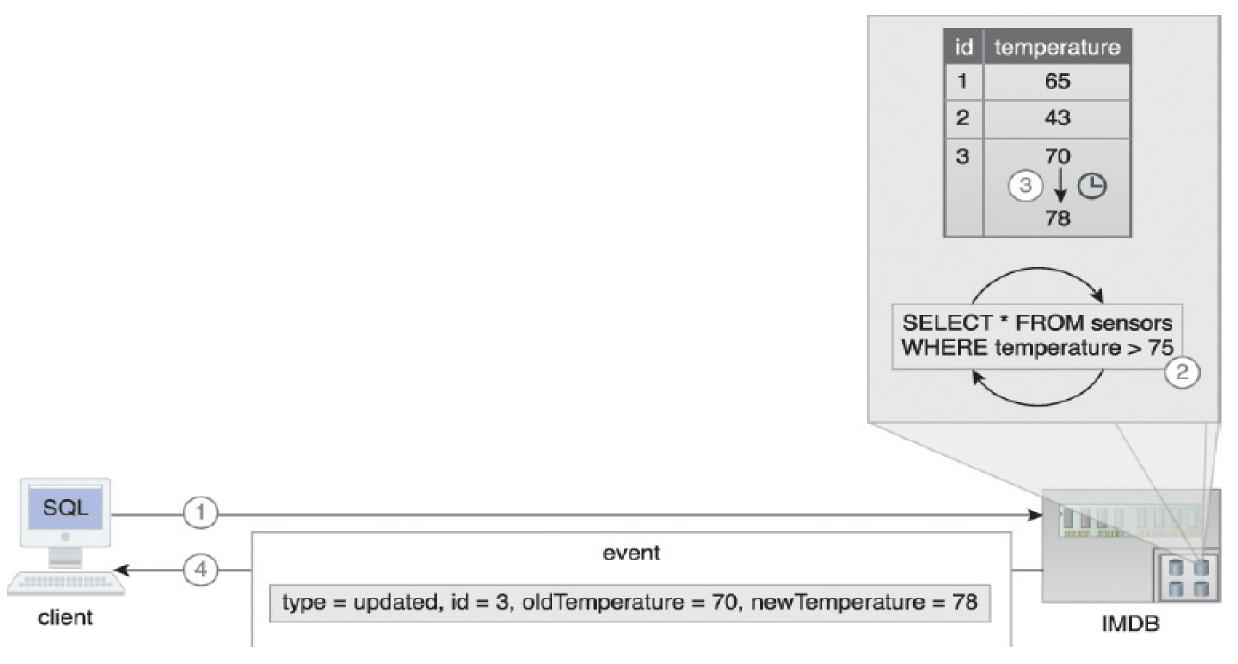
In-Memory Storage: In-Memory Database

- ➤IMDBs are in-memory storage devices that **employ database technology** and leverage the performance of **RAM** to **overcome runtime latency** issues of the on-disk storage devices.
- An IMDB can be relational in nature (relational IMDB) for the storage of structured data, or may leverage NoSQL technology (non-relational IMDB) for the storage of semistructured and unstructured data.



An example depicting the retrieval of data from an IMDB.

- The previous example describes the following scenario:
 - 1. A relational dataset is **stored into an IMDB**.
 - 2. A **client requests** a customer record (id = 2) via SQL.
 - 3. The relevant customer **record is then returned** by the **IMDB**, which is directly manipulated by the client without the need for any deserialization.



An example of IMDB storage configured with a continuous query.

- The previous example describes the following scenario:
 - 1. A client issues a **continuous query** (select * from sensors where temperature > 75).
 - 2. It is registered in the **IMDB**.
 - 3. When the temperature for any sensor exceeds 75F ...
 - 4. ... an **updated event** is sent to the subscribing **client** that contains various details about the event.

- ➤IMDBs are heavily used in **real-time analytics** and can further be used for developing **low latency applications** requiring full ACID transaction support (relational IMDB).
- Introduction of IMDBs into an existing Big Data solution generally requires **replacement of existing on-disk storage devices**, including any RDBMSs if used.
- Examples include Aerospike, MemSQL, Altibase HDB, eXtreme DB and Pivotal GemFire XD

- In the case of replacing an RDBMS with a relational IMDB, **little or no** application code change is required due to SQL support provided by the relational IMDB.
- However, when replacing an RDBMS with a NoSQL IMDB, code change may be required due to the need to implement the IMDB's NoSQL APIs.
- In the case of replacing an on-disk NoSQL database with a relational IMDB, code change will often be required to establish SQL-based access.
- However, when replacing an on-disk NoSQL database with a NoSQL IMDB, code change may still be required due to the implementation of new APIs.

Thank You